

# Format String Syntax

Formatting functions such as `fmt::format()` ([api.html#format](#)) and `fmt::print()` ([api.html#print](#)) use the same format string syntax described in this section.

Format strings contain “replacement fields” surrounded by curly braces `{}`. Anything that is not contained in braces is considered literal text, which is copied unchanged to the output. If you need to include a brace character in the literal text, it can be escaped by doubling: `{{` and `}}`.

The grammar for a replacement field is as follows:

```
replacement_field ::= "{" [arg_id] [":" format_spec] "}"
arg_id            ::= integer | identifier
integer           ::= digit+
digit             ::= "0"..."9"
identifier        ::= id_start id_continue*
id_start          ::= "a"..."z" | "A"..."Z" | "_"
id_continue       ::= id_start | digit
```

In less formal terms, the replacement field can start with an *arg\_id* that specifies the argument whose value is to be formatted and inserted into the output instead of the replacement field. The *arg\_id* is optionally followed by a *format\_spec*, which is preceded by a colon `:`. These specify a non-default format for the replacement value.

See also the [Format Specification Mini-Language](#) section.

If the numerical *arg\_ids* in a format string are 0, 1, 2, ... in sequence, they can all be omitted (not just some) and the numbers 0, 1, 2, ... will be automatically inserted in that order.

Named arguments can be referred to by their names or indices.

Some simple format string examples:

```
"First, thou shalt count to {0}" // References the first argument
"Bring me a {}"                 // Implicitly references the first argument
"From {} to {}".format(1, 2)    // Same as "From {0} to {1}"
```

The *format\_spec* field contains a specification of how the value should be presented, including such details as field width, alignment, padding, decimal precision and so on. Each value type can define its own “formatting mini-language” or interpretation of the *format\_spec*.

Most built-in types support a common formatting mini-language, which is described in the next section.

A *format\_spec* field can also include nested replacement fields in certain positions within it. These nested replacement fields can contain only an argument id; format specifications are not allowed. This allows the formatting of a value to be dynamically specified.

See the [Format Examples](#) section for some examples.

## Format Specification Mini-Language

“Format specifications” are used within replacement fields contained within a format string to define how individual values are presented (see [Format String Syntax](#)). Each formattable type may define how the format specification is to be interpreted.

Most built-in types implement the following options for format specifications, although some of the formatting options are only supported by the numeric types.

The general form of a *standard format specifier* is:

```
format_spec ::= [[fill]align][sign]["#"]["0"][width]["." precision][type]
fill        ::= <a character other than '{' or '}'>
align       ::= "<" | ">" | "^"
sign        ::= "+" | "-" | " "
width       ::= integer | "{" [arg_id] "}"
precision   ::= integer | "{" [arg_id] "}"
type        ::= int_type | "a" | "A" | "c" | "e" | "E" | "f" | "F" | "g" | "G" | "L" | "p" | "r"
int_type    ::= "b" | "B" | "d" | "o" | "x" | "X"
```

The *fill* character can be any Unicode code point other than `'{'` or `'}'`. The presence of a fill character is signaled by the character following it, which must be one of the alignment options. If the second character of *format\_spec* is not a valid alignment option, then it is assumed that both the fill character and the alignment option are absent.

The meaning of the various alignment options is as follows:

Option	Meaning
'<'	Forces the field to be left-aligned within the available space (this is the default for most objects).
'>'	Forces the field to be right-aligned within the available space (this is the default for numbers).
'^'	Forces the field to be centered within the available space.

Note that unless a minimum field width is defined, the field width will always be the same size as the data to fill it, so that the alignment option has no meaning in this case.

The *sign* option is only valid for number types, and can be one of the following:

Option	Meaning
'+'	indicates that a sign should be used for both positive as well as negative numbers.
'-'	indicates that a sign should be used only for negative numbers (this is the default behavior).
space	indicates that a leading space should be used on positive numbers, and a minus sign on negative numbers.

The `'#'` option causes the “alternate form” to be used for the conversion. The alternate form is defined differently for different types. This option is only valid for integer and floating-point types. For integers, when binary, octal, or hexadecimal output is used, this option adds the prefix respective `"0b"` (`"0B"`), `"0o"`, or `"0x"` (`"0X"`) to the output value. Whether the prefix is lower-case or upper-case is determined by the case of the type specifier, for example, the prefix `"0x"` is used for the type `'x'` and `"0X"` is used for `'X'`. For floating-point numbers the alternate form causes the result of the conversion to always contain a decimal-point character, even if no digits follow it. Normally, a decimal-point character appears in the result of these conversions only if a digit follows it. In addition, for `'g'` and `'G'` conversions, trailing zeros are not removed from the result.

*width* is a decimal integer defining the minimum field width. If not specified, then the field width will be determined by the content.

Preceding the *width* field by a zero (`'0'`) character enables sign-aware zero-padding for numeric types. It forces the padding to be placed after the sign or base (if any) but before the digits. This is used for printing fields in the form `'+000000120'`. This option is only valid for numeric types and it has no effect on formatting of infinity and NaN.

The *precision* is a decimal number indicating how many digits should be displayed after the decimal point for a floating-point value formatted with 'f' and 'F', or before and after the decimal point for a floating-point value formatted with 'g' or 'G'. For non-number types the field indicates the maximum field size - in other words, how many characters will be used from the field content. The *precision* is not allowed for integer, character, Boolean, and pointer values.

Finally, the *type* determines how the data should be presented.

The available string presentation types are:

Type	Meaning
's'	String format. This is the default type for strings and may be omitted.
none	The same as 's'.

The available character presentation types are:

Type	Meaning
'c'	Character format. This is the default type for characters and may be omitted.
none	The same as 'c'.

The available integer presentation types are:

Type	Meaning
'b'	Binary format. Outputs the number in base 2. Using the '#' option with this type adds the prefix "0b" to the output value.
'B'	Binary format. Outputs the number in base 2. Using the '#' option with this type adds the prefix "0B" to the output value.
'd'	Decimal integer. Outputs the number in base 10.
'o'	Octal format. Outputs the number in base 8.
'x'	Hex format. Outputs the number in base 16, using lower-case letters for the digits above 9. Using the '#' option with this type adds the prefix "0x" to the output value.
'X'	Hex format. Outputs the number in base 16, using upper-case letters for the digits above 9. Using the '#' option with this type adds the prefix "0X" to the output value.
'L'	Locale-specific format. This is the same as 'd', except that it uses the current locale setting to insert the appropriate number separator characters.
none	The same as 'd'.

Integer presentation types can also be used with character and Boolean values. Boolean values are formatted using textual representation, either `true` or `false`, if the presentation type is not specified.

The available presentation types for floating-point values are:

Type	Meaning
'a'	Hexadecimal floating point format. Prints the number in base 16 with prefix "0x" and lower-case letters for digits above 9. Uses 'p' to indicate the exponent.
'A'	Same as 'a' except it uses upper-case letters for the prefix, digits above 9 and to indicate the exponent.
'e'	Exponent notation. Prints the number in scientific notation using the letter 'e' to indicate the exponent.
'E'	Exponent notation. Same as 'e' except it uses an upper-case 'E' as the separator character.
'f'	Fixed point. Displays the number as a fixed-point number.
'F'	Fixed point. Same as 'f', but converts <code>nan</code> to <code>NAN</code> and <code>inf</code> to <code>INF</code> .

Type	Meaning
'g'	General format. For a given precision $p \geq 1$ , this rounds the number to $p$ significant digits and then formats the result in either fixed-point format or in scientific notation, depending on its magnitude.  A precision of $0$ is treated as equivalent to a precision of $1$ .
'G'	General format. Same as 'g' except switches to 'E' if the number gets too large. The representations of infinity and NaN are uppercased, too.
'L'	Locale-specific format. This is the same as 'g', except that it uses the current locale setting to insert the appropriate number separator characters.
none	Similar to 'g', except that fixed-point notation, when used, has at least one digit past the decimal point. The default precision is as high as needed to represent the particular value.

The available presentation types for pointers are:

Type	Meaning
'p'	Pointer format. This is the default type for pointers and may be omitted.
none	The same as 'p'.

## Format Examples

This section contains examples of the format syntax and comparison with the printf formatting.

In most of the cases the syntax is similar to the printf formatting, with the addition of the `{}` and with `:` used instead of `%`. For example, `"%03.2f"` can be translated to `"{:03.2f}"`.

The new format syntax also supports new and different options, shown in the following examples.

Accessing arguments by position:

```
fmt::format("{0}, {1}, {2}", 'a', 'b', 'c');
// Result: "a, b, c"
fmt::format("{}, {}, {}", 'a', 'b', 'c');
// Result: "a, b, c"
fmt::format("{2}, {1}, {0}", 'a', 'b', 'c');
// Result: "c, b, a"
fmt::format("{0}{1}{0}", "abra", "cad"); // arguments' indices can be repeated
// Result: "abracadabra"
```

Aligning the text and specifying a width:

```
fmt::format("{:<30}", "left aligned");
// Result: "left aligned"
fmt::format("{:>30}", "right aligned");
// Result: "right aligned"
fmt::format("{:^30}", "centered");
// Result: "centered"
fmt::format("{:*^30}", "centered"); // use '*' as a fill char
// Result: "*****centered*****"
```

Dynamic width:

```
fmt::format("{:<{}}", "left aligned", 30);
// Result: "left aligned"
```

Dynamic precision:

```
fmt::format("{:.{}}f", 3.14, 1);  
// Result: "3.1"
```

Replacing %f , %-f , and % f and specifying a sign:

```
fmt::format("{:+f}; {:+f}", 3.14, -3.14); // show it always  
// Result: "+3.140000; -3.140000"  
fmt::format("{: f}; {: f}", 3.14, -3.14); // show a space for positive numbers  
// Result: " 3.140000; -3.140000"  
fmt::format("{:-f}; {:-f}", 3.14, -3.14); // show only the minus -- same as '{:f}; {:f}'  
// Result: "3.140000; -3.140000"
```

Replacing %x and %o and converting the value to different bases:

```
fmt::format("int: {0:d}; hex: {0:x}; oct: {0:o}; bin: {0:b}", 42);  
// Result: "int: 42; hex: 2a; oct: 52; bin: 101010"  
// with 0x or 0 or 0b as prefix:  
fmt::format("int: {0:d}; hex: {0:#x}; oct: {0:#o}; bin: {0:#b}", 42);  
// Result: "int: 42; hex: 0x2a; oct: 052; bin: 0b101010"
```

Padded hex byte with prefix and always prints both hex characters:

```
fmt::format("{:#04x}", 0);  
// Result: "0x00"
```

Box drawing using Unicode fill:

```
fmt::print(  
    "┌{0:-^{2}}┐\n"  
    "│{1: ^{2}}│\n"  
    "└{0:-^{2}}┘\n", "", "Hello, world!", 20);
```

prints:

```
┌ Hello, world! ┐
```

Using type-specific formatting:

```
#include <fmt/chrono.h>  
  
auto t = tm();  
t.tm_year = 2010 - 1900;  
t.tm_mon = 6;  
t.tm_mday = 4;  
t.tm_hour = 12;  
t.tm_min = 15;  
t.tm_sec = 58;  
fmt::print("{:%Y-%m-%d %H:%M:%S}", t);  
// Prints: 2010-08-04 12:15:58
```

Using the comma as a thousands separator:

```
#include <fmt/locale.h>
```

```
auto s = fmt::format(std::locale("en_US.UTF-8"), "{:L}", 1234567890);
```

```
// s == "1,234,567,890"
```